

# Operating Systems

Jinghui Zhong (钟竞辉)


Office: B3-515

Email : [jinghuizhong@scut.edu.cn](mailto:jinghuizhong@scut.edu.cn)




# Inter Process Communication (IPC)

- How to pass information among processes?
- How to make sure two or more processes do not get into each other's way when engaging in critical activities.



- Proper sequencing when dependencies are present.



# Race Conditions

- **Race conditions:** situations in which several processes access shared data and the final result depends on the order of operations.
- With increasing parallelism due to increasing number of cores, race condition are becoming more common.

# Example of Race Condition

**Out:** points to the next file to be printed

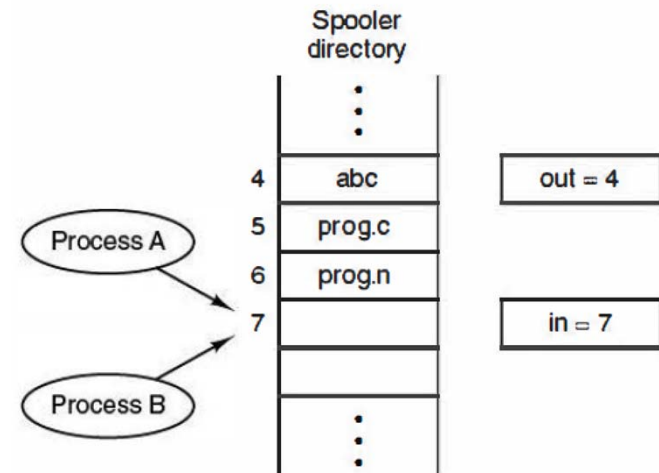
**In:** points to the next free slot in the directory.

**in** = 7

(1): Process A reads **in** and stores the value 7 in a local variable. **A switch to process B happens.**

(2): Process B reads **in**, stores the file name in slot 7 and updates **in** to be an 8.

(3): Process A stores the file name in slot 7 and updates **in** to be an 8.



The file name in slot 7 was determined by who finished last. A race condition occurs.



# Critical Regions

- Key idea to avoid race condition: **prohibit more than one process from reading and writing the shared data at the same time.**
- **Critical Region:** part of the program where the share memory is accessed.

```
// 临界区结构对象
CRITICAL_SECTION g_cs;
// 共享资源
char g_cArray[10];
UINT ThreadProc10(LPVOID pParam)
{
    // 进入临界区
    EnterCriticalSection(&g_cs);
    // 对共享资源进行写入操作
    for (int i = 0; i < 10; i++)
    {
        g_cArray[i] = a;
        Sleep(1);
    }
    // 离开临界区
    LeaveCriticalSection(&g_cs);
    return 0;
}
UINT ThreadProc11(LPVOID pParam)
{
    // 进入临界区
    EnterCriticalSection(&g_cs);
    // 对共享资源进行写入操作
    for (int i = 0; i < 10; i++)
    {
        g_cArray[10 - i - 1] = b;
        Sleep(1);
    }
    // 离开临界区
    LeaveCriticalSection(&g_cs);
    return 0;
}
.....
void CSample08View::OnCriticalSection()
{
    // 初始化临界区
    InitializeCriticalSection(&g_cs);
    // 启动线程
    AfxBeginThread(ThreadProc10, NULL);
    AfxBeginThread(ThreadProc11, NULL);
    // 等待计算完毕
    Sleep(300);
    // 报告计算结果
    CString sResult = CString(g_cArray);
    AfxMessageBox(sResult);
}
```

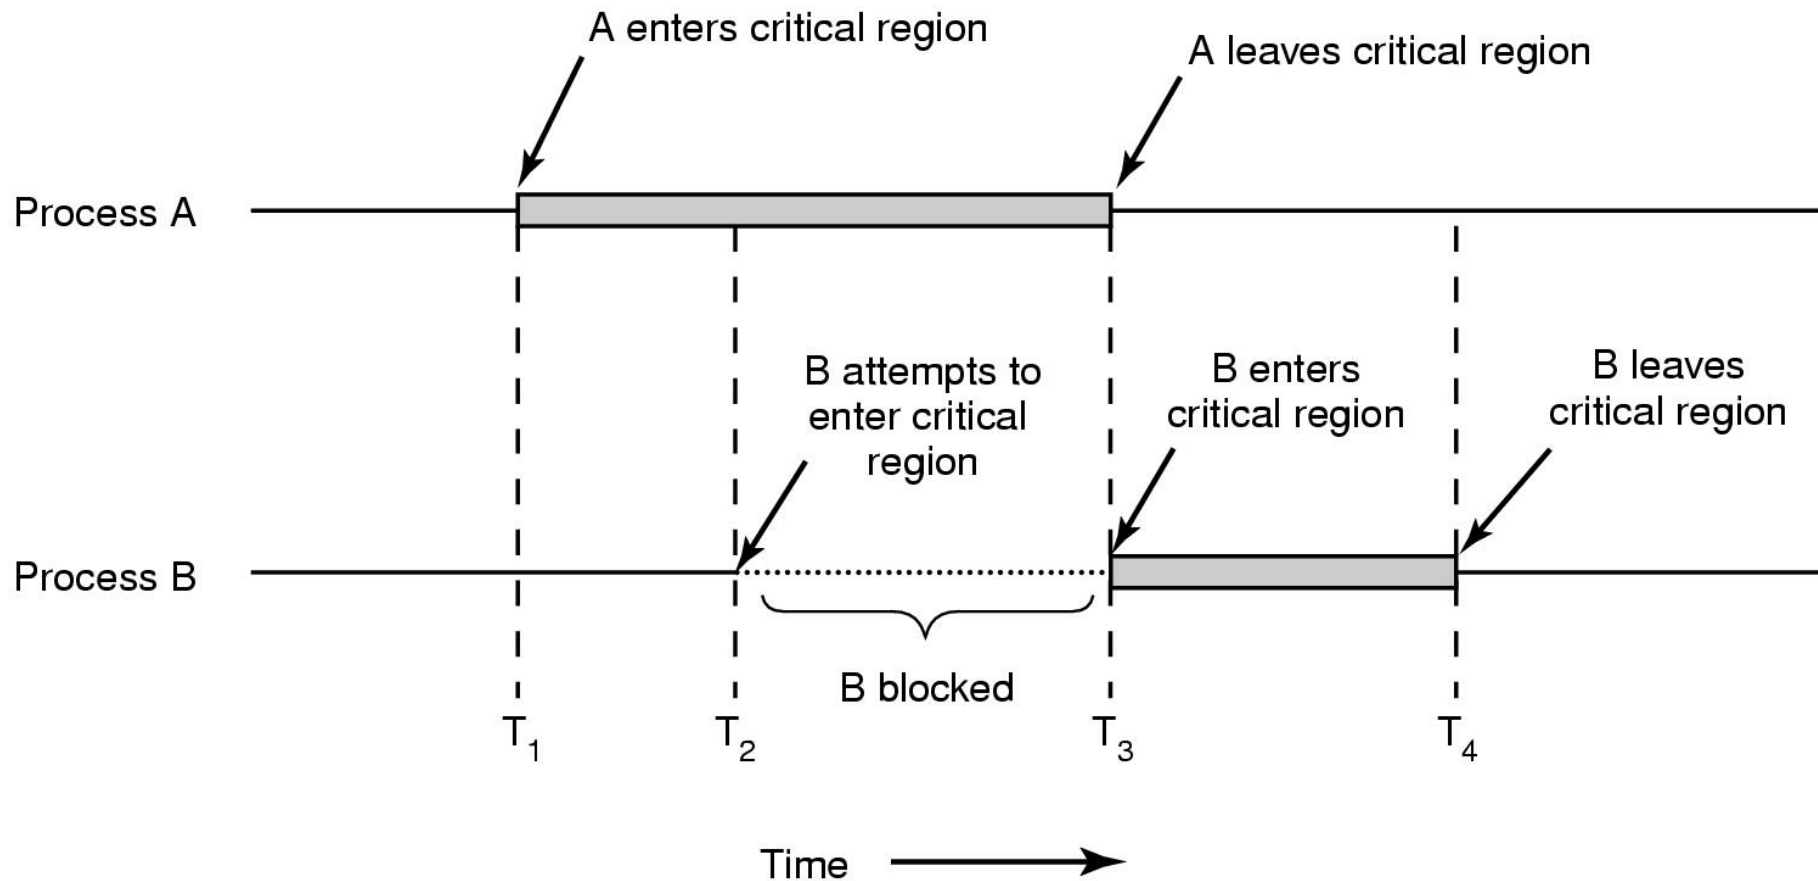


# Critical Regions

## ● Four conditions to support a good solution

- ① No two processes may be simultaneously in critical region
- ② No assumption made about speeds or numbers of CPUs
- ③ No process running outside its critical region may block another process
- ④ No process must wait forever to enter its critical region

# Mutual Exclusion using Critical Regions



# Mutual Exclusion Solution

## - Disabling Interrupts

- The CPU is only switch from process to process when clock or other interrupts happen; Hence, by disabling all interrupts, the CPU will not be switched to another process.
- However, it is unwise to allow user processes to disable interrupts.
  - ✓ One thread may never turn on interrupt;
  - ✓ Problem still exist for multiprocessor systems.



# Mutual Exclusion Solution - Lock Variable

```
shared int lock = 0;
```

```
/* entry_code: execute before entering critical section */
```

```
while (lock != 0) ; // do nothing
```

```
lock = 1;
```

- critical section -

```
/* exit_code: execute after leaving critical section */
```

```
lock = 0;
```

## Problem?

- If a context switch occurs after one process executes the while statement, but before setting lock = 1, then two (or more) processes may be able to enter their critical sections at the same time.

# Mutual Exclusion Solution

## – Strict Alternation

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region( );  
}
```

(a) Process 0.

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region( );  
}
```

(b) Process 1.

- Since the processes must strictly alternate entering their critical sections, a process wanting to enter its critical section twice will be blocked until the other process decides to enter (and leave) its critical section.

# Mutual Exclusion – Peterson's Solution

```
#define FALSE 0
#define TRUE 1
#define N      2          /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;              /* number of the other process */

    other = 1 - process;    /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;         /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

● This solution satisfies all 4 properties of a good solution. Unfortunately, this solution involves **busy waiting** in the while loop.



# Hardware solution: Test-and-Set Locks (TSL)

- The hardware must support a special instruction, TSL, which does **two** things in a single atomic action:
  - (a) copy a value in memory (flag) to a CPU register
  - (b) set flag to 1.

enter\_region:

TSL REGISTER, LOCK	copy lock to register and set lock to 1
CMP REGISTER, #0	was lock zero?
JNE enter_region	if it was non zero, lock was set, so loop
RET	return to caller; critical region entered

leave\_region:

MOVE LOCK, #0	store a 0 in lock
RET	return to caller



# Mutual Exclusion with Busy Waiting

- **BUSY-WAITING**: a process executing the entry code will sit in a tight loop using up CPU cycles, testing some condition over and over, until it becomes true.
- Busy-waiting may lead to the **priority-inversion problem**.

Not only does this approach waste CPU time, but it can also have unexpected effects. Consider a computer with two processes,  $H$ , with high priority, and  $L$ , with low priority. The scheduling rules are such that  $H$  is run whenever it is in ready state. At a certain moment, with  $L$  in its critical region,  $H$  becomes ready to run (e.g., an I/O operation completes).  $H$  now begins busy waiting, but since  $L$  is never scheduled while  $H$  is running,  $L$  never gets the chance to leave its critical region, so  $H$  loops forever. This situation is sometimes referred to as the **priority inversion problem**.



# Sleep and Wakeup

- From busy waiting to blocking...


**Sleep:** is a system call that causes the caller to block, that is, be suspended until another process wakes it up.

**Wakeup:** has one parameter, the process to be awakened.



# Producer-Consumer Problem

- **Producer-Consumer Problem:** Consider a circular buffer that can hold  $N$  items. Producers add items to the buffer and Consumers remove items from the buffer. The Producer-Consumer Problem is to restrict access to the buffer.



# Producer-Consumer Problem

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                   /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        item = produce_item();                    /* generate next item */
        if (count == N) sleep();                  /* if buffer is full, go to sleep */
        insert_item(item);                        /* put item in buffer */
        count = count + 1;                        /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);        /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        if (count == 0) sleep();                  /* if buffer is empty, got to sleep */
        item = remove_item();                    /* take item out of buffer */
        count = count - 1;                        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);    /* was buffer full? */
        consume_item(item);                      /* print item */
    }
}
```

Producer-consumer problem with fatal race condition



# Semaphores [E.W. Dijkstra, 1965]

- A SEMAPHORE, S, is a structure consisting of two parts:

- (a) an integer counter, COUNT

- (b) a queue of pids of blocked processes, Q

- That is,

```
struct sem_struct {  
    int count;  
    queue Q;  
} semaphore;
```

```
semaphore S;
```



# Semaphores [E.W. Dijkstra, 1965]

- There are two operations on semaphores, **UP** and **DOWN** (PV). These operations must be **executed atomically** (that is in mutual exclusion). Suppose that P is the process making the system call. The operations are defined as follows:

DOWN(S):

if (S.count > 0)

    S.count = S.count - 1;

else

    block(P); that is,

    (a) enqueue the pid of P in S.Q,

    (b) block process P (remove the pid from the ready queue)

    (c) pass control to the scheduler.



# Semaphores [E.W. Dijkstra, 1965]

UP(S):

if (S.Q is nonempty)

wakeup(P) for some process P in S.Q; that is,

(a) remove a pid from S.Q (the pid of P),

(b) put the pid in the ready queue, and

(c) pass control to the scheduler.

else

S.count = S.count + 1;

# Mutual Exclusion Solution

- Semaphores do not require busy-waiting, instead they involve BLOCKING.

```
semaphore mutex = 1; // set mutex.count = 1
```

```
DOWN(mutex);
```

```
- critical section -
```

```
UP(mutex);
```



# Mutexes

- A mutex is a semaphore that can be in one of two states: unlocked (0) or locked (1).

mutex\_lock:

TSL REGISTER,MUTEX

| copy mutex to register and set mutex to 1

CMP REGISTER,#0

| was mutex zero?

JZE ok

| if it was zero, mutex was unlocked, so return

CALL thread\_yield

| mutex is busy; schedule another thread

JMP mutex\_lock

| try again later

ok: RET | return to caller; critical region entered

mutex\_unlock:

MOVE MUTEX,#0

| store a 0 in mutex

RET | return to caller

Implementation of *mutex\_lock* and *mutex\_unlock*



# Previous Solution

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                   /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                               /* repeat forever */
        item = produce_item();                  /* generate next item */
        if (count == N) sleep();                /* if buffer is full, go to sleep */
        insert_item(item);                      /* put item in buffer */
        count = count + 1;                      /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);      /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                               /* repeat forever */
        if (count == 0) sleep();                /* if buffer is empty, got to sleep */
        item = remove_item();                  /* take item out of buffer */
        count = count - 1;                     /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                   /* print item */
    }
}
```

Producer-consumer problem with fatal race condition



# Semaphore Based Solution

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```


The producer-consumer problem using semaphores



# Using Semaphores

## ● Process Synchronization:

Suppose we have 4 processes: A, B, C, and D. A must finish executing before B start. B and C must finish executing before D starts.



**How many semaphores should be used to achieve the above goal?**

# Using Semaphores

## ● Process Synchronization:

Process A:

- do work of A

UP(S1);      /\* Let B or C start \*/

Process B:

DOWN(S1);   /\* Block until A is finished \*/

- do work of B

UP(S2);

Process C:

- do work of C


UP(S3);

Process D:

DOWN(S2);

DOWN(S3);

- do work of D



# Check Points

- What is Race Condition?
- What is Critical Region?
- What is Busy Waiting?
- What is Semaphore?
- What is Mutex?