

# Operating Systems

Jinghui Zhong (钟竞辉)

Office: B3-515

Email : [jinghuizhong@scut.edu.cn](mailto:jinghuizhong@scut.edu.cn)

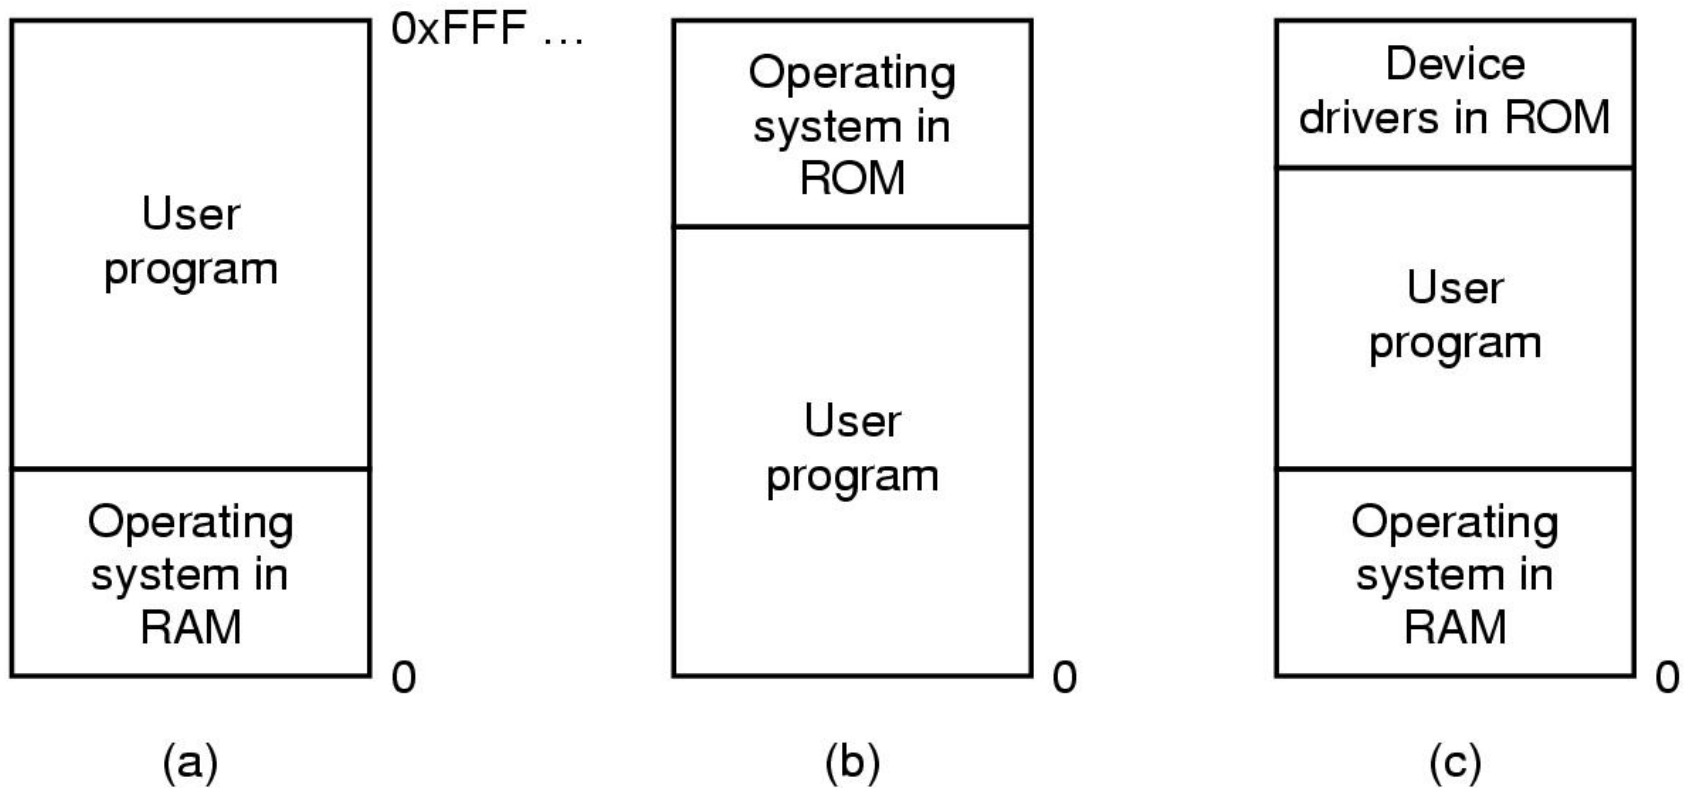


# Memory Management

- Ideally programmers want memory that is
  - ✓ large
  - ✓ fast
  - ✓ non volatile
- Memory hierarchy
  - ✓ small amount of fast, expensive memory – cache
  - ✓ some medium-speed, medium price main memory
  - ✓ gigabytes of slow, cheap disk storage
- Memory manager handles the memory hierarchy

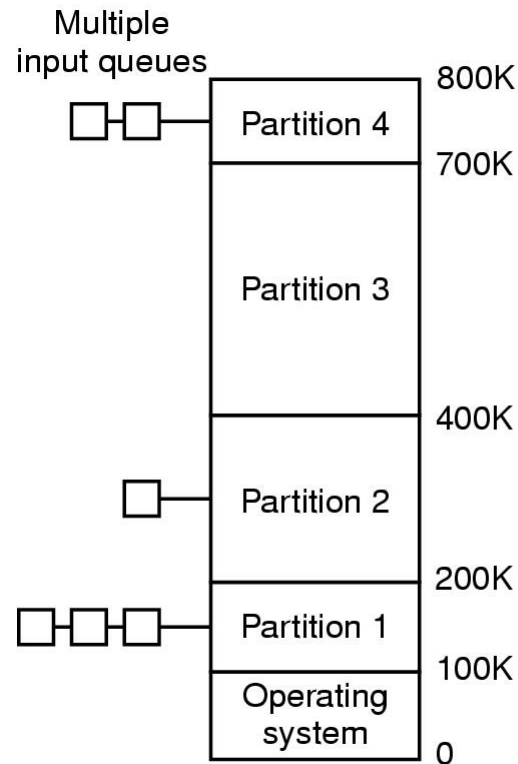


# Basic Memory Management




Three simple ways of organizing memory  
- an operating system with one user process

# Multiprogramming with Fixed Partitions



(a)



(b)


- Separate input queues for each partition
- Single input queue
- As the partition sizes are fixed, any space not used by a particular job is lost.
- It may not be easy to state how big a partition a particular job needs.

# Swapping & Virtual Memory

- Two approaches to overcome the limitation of memory:
  - ✓ **Swapping** puts a process back and forth in memory and on the disk.
  - ✓ **Virtual memory** allows programs to run even when they are only partially in main memory.



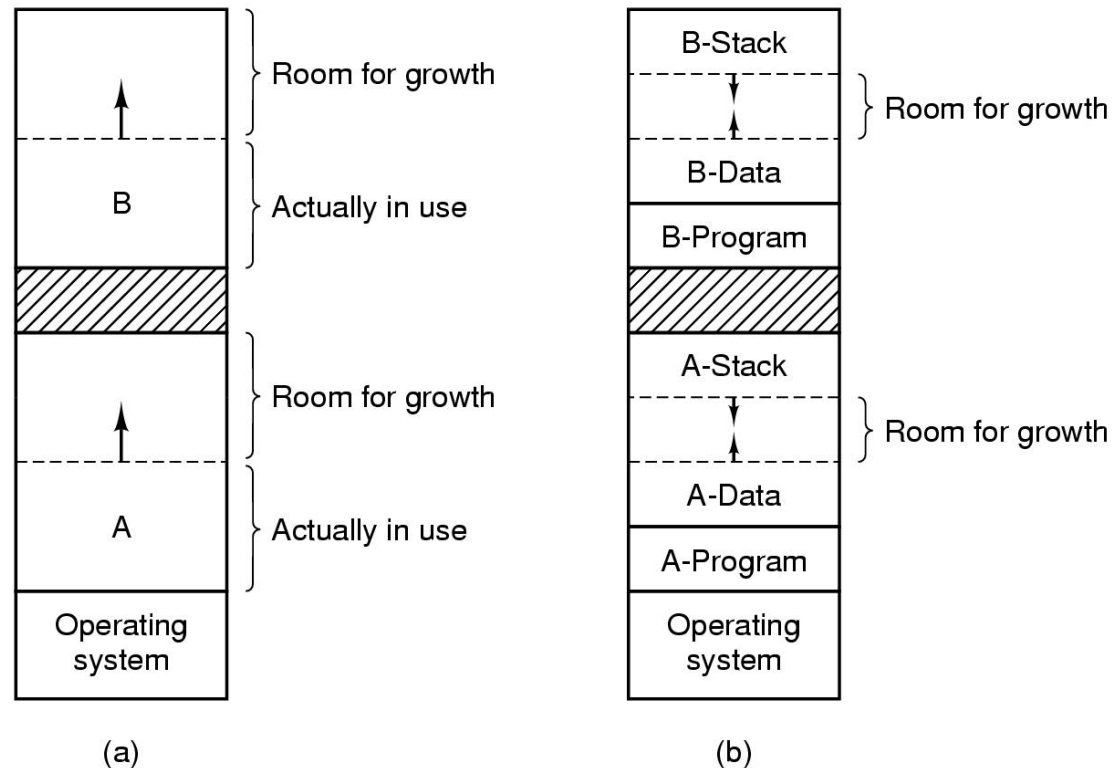
# Swapping



Memory allocation changes as processes come into memory and leave memory.

- When swapping creates multiple holes in memory, **memory compaction** can be used to combine them into a big one by moving all processes together.


# Swapping



- Allocating space for growing data segment
- Allocating space for growing stack & data segment

# Memory Management with Bit Maps & List

- Drawback of bitmaps: to find consecutive 0 bits in the map is time-consuming.



(a) Part of memory with 5 processes, 3 holes

- ✓ tick marks show allocation units
- ✓ shaded regions are free


(b) Corresponding bit map

(c) Same information as a list



# Memory Management with Linked Lists

## ● Four neighbor combinations for the terminating process



## ● Four algorithms for memory management:


- ✓ **First fit:** Searches from the beginning for a hole that fits.
- ✓ **Next fit:** Searches from the place where it left off last time for a hole that fits.
- ✓ **Best fit:** Searches the entire list and takes the smallest hole that fits.
- ✓ **Worst fit:** Searches the largest hole that fits.

# Virtual Memory

- Problem: Program too large to fit in memory
- Solutions:
  - ✓ Programmer splits program into pieces called **Overlays** - too much work
  - ✓ **Virtual memory** - OS keeps the part of the program currently in use in memory
- **Paging** is a technique used to implement virtual memory.
- **Virtual Address** is a program generated address.
- The **MMU** (memory management unit) translates a virtual address into a physical address.



# MMU



The position and function of the MMU

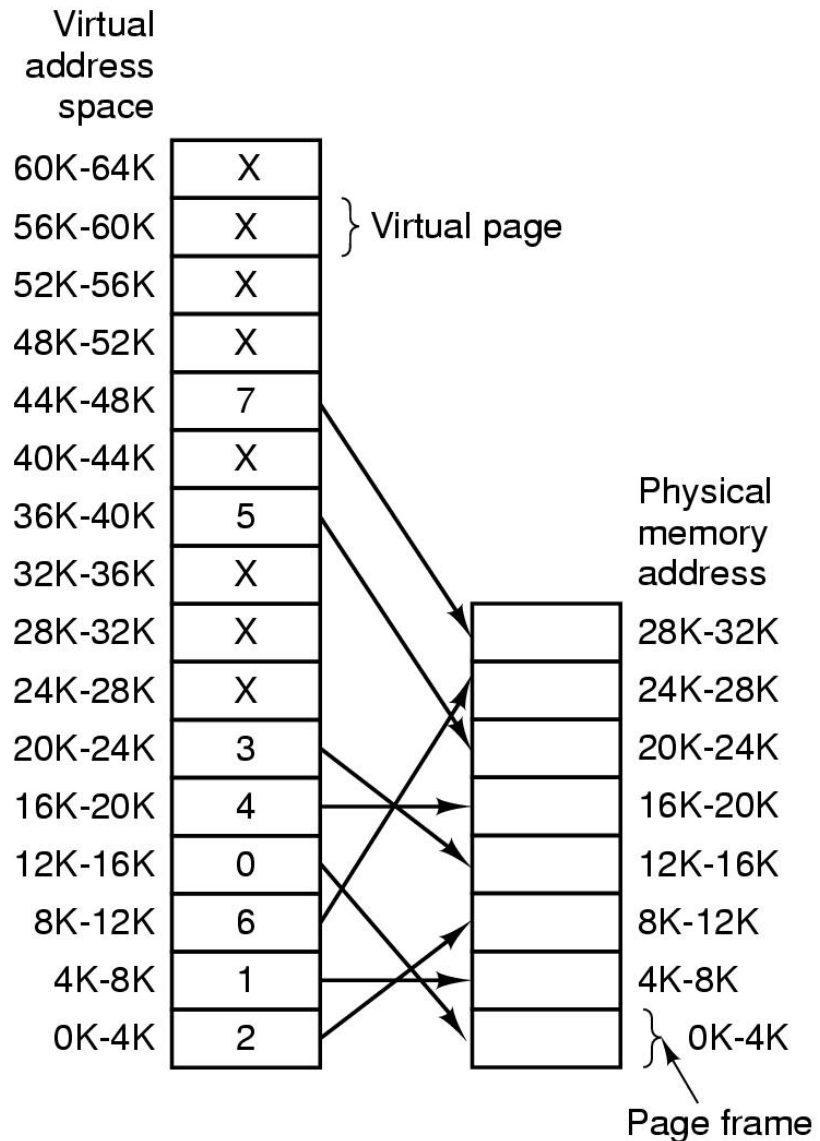
# Virtual Memory

- The virtual address space is divided into (virtual) **pages** and those in the physical memory are (page) **frames**.
- A Present/Absent bit keeps track of whether or not the page is mapped.
- Reference to an unmapped page causes the CPU to trap to the OS.
- This trap is called a **Page fault**. The MMU selects a little used page frame, writes its contents back to disk, fetches the page just referenced, and restarts the trapped instruction.



# Page Table

The relation between virtual addresses and physical memory addresses given by page table



# Page Tables

- Example: Virtual address = 4097 = 0001 000000000001


Virtual page #   12-bit offset

- The purpose of the page table is to map virtual pages into page frames. The page table is a function to map the virtual page to the page frame.
- Two major issues :
  - ✓ Page tables may be extremely large (e.g. most computers use) 32-bit address with 4k page size, 12-bit offset
    - ➔ 20 bits for virtual page number
    - ➔ 1 million entries!
  - ✓ The mapping must be fast because it is done on every memory access!!






# Pure paging



# Page Table

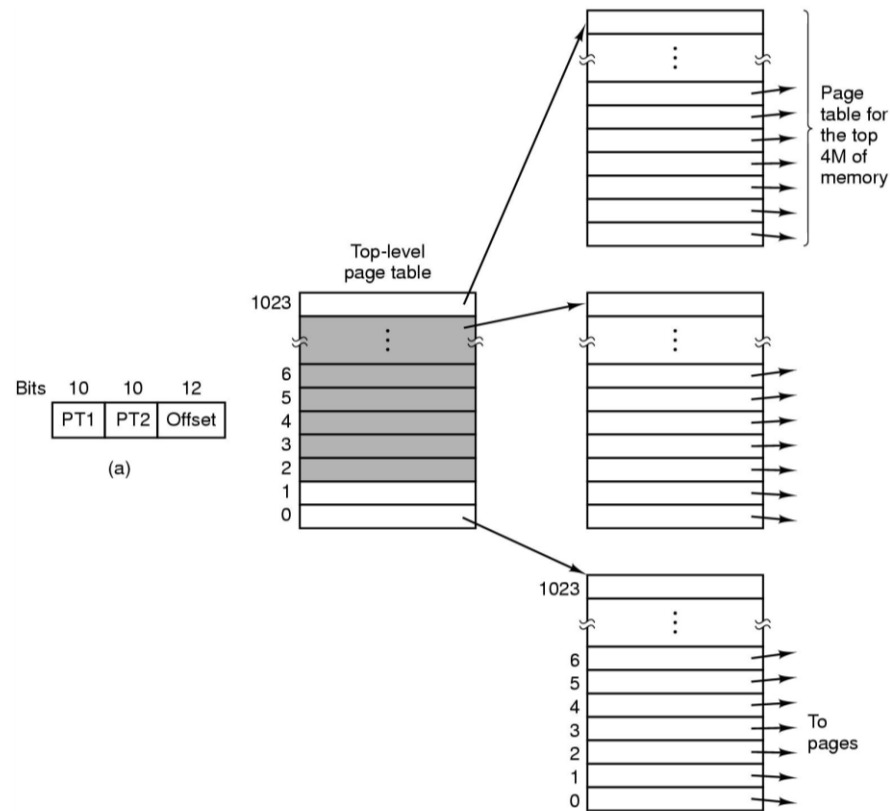


Internal operation of MMU with 16 4 KB pages




# Multilevel Page Tables

- Multilevel page tables - **reduce the table size**. Also, **don't keep page tables in memory that are not needed**.



32 bit address with 2 page table fields

# Two-Level Page-Table Scheme



# Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
  - ✓ a page number consisting of 20 bits.
  - ✓ a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
  - ✓ a 10-bit page number.
  - ✓ a 10-bit page offset.
- Thus, a logical address is as follows:

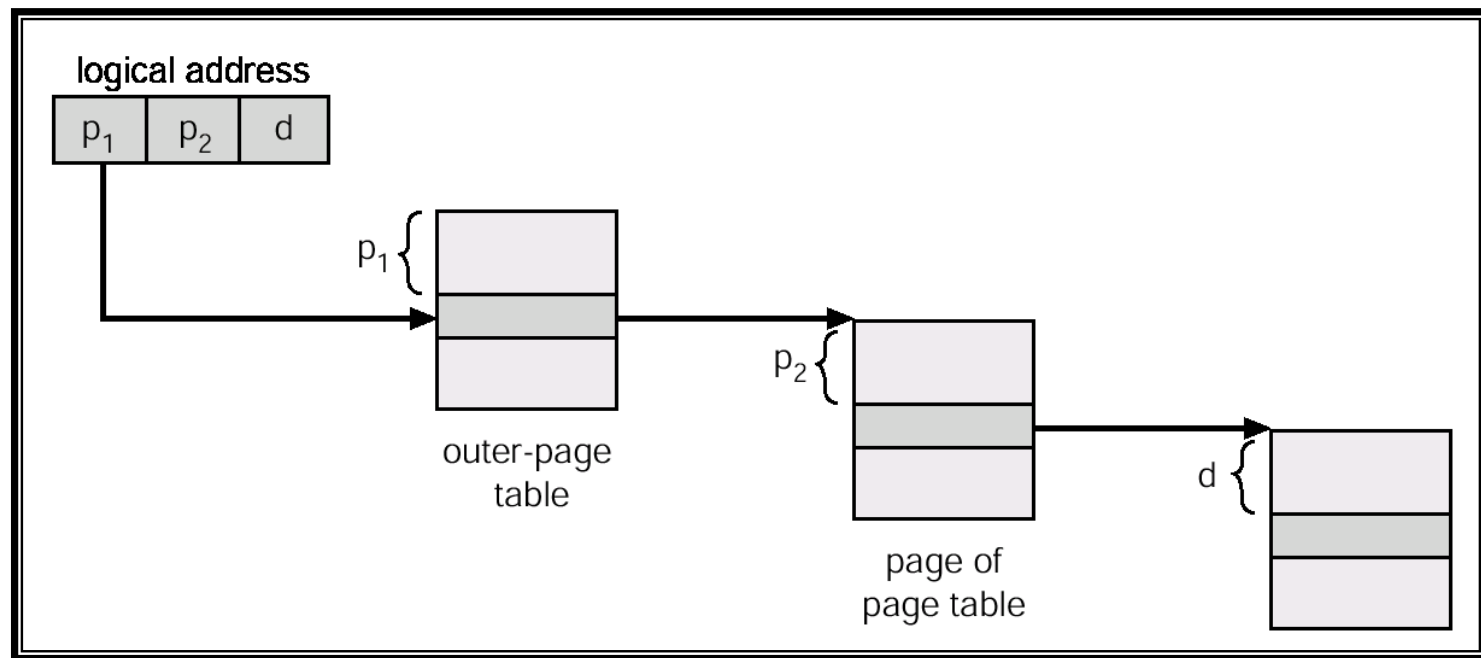
page number		page offset
$p_1$	$p_2$	$d$
10	10	12

where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table.



# Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture is shown as below.

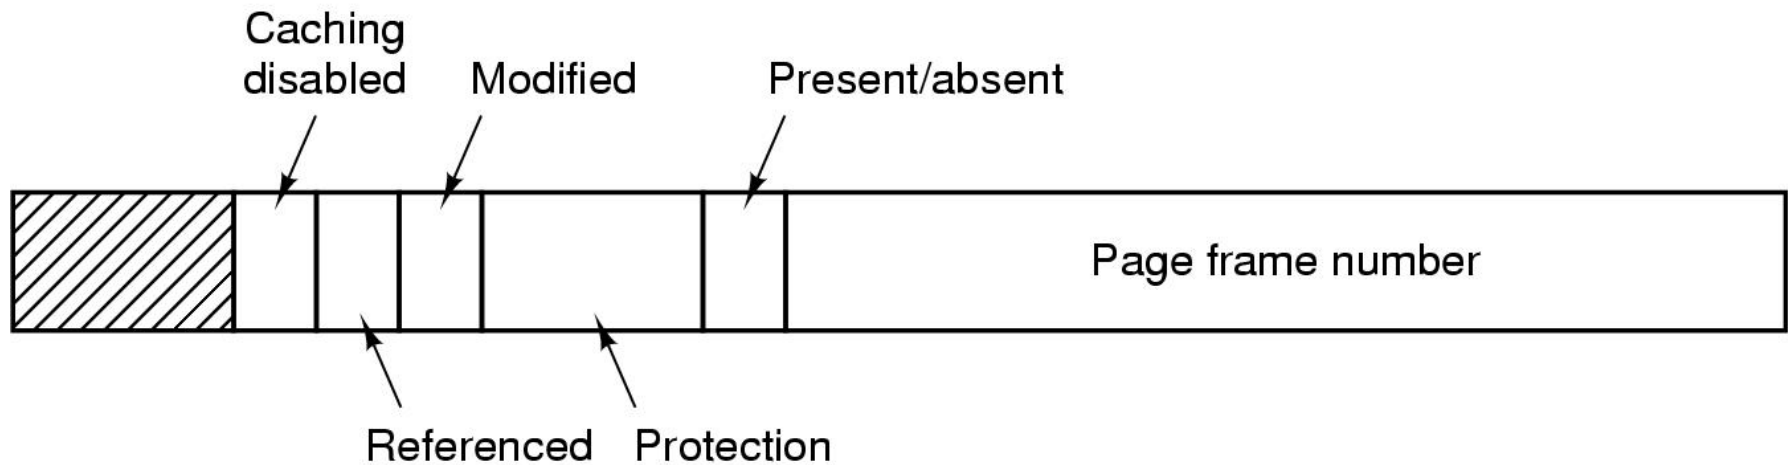




# Page Tables

- Most OSs allocate a page table for each process.
- Single page table consisting of an array of hardware registers. As a process is loaded, the registers are loaded with page table.
  - ✓ Advantage - simple
  - ✓ Disadvantage - expensive if table is large and loading the full page table at every context switch hurts performance.
- Leave page table in memory - a single register points to the table
  - ✓ Advantage - context switch cheap
  - ✓ Disadvantage - one or more memory references to read table entries

# Page Tables



Typical page table entry

# Structure of a Page Table Entry

- Page frame number: map the frame number
- Present/absent bit: 1/0 indicates valid/invalid entry
- Protection bit: what kinds of access are permitted.
- Modified – set when modified and writing to the disk occur
- Referenced - Set when page is referenced (help decide which page to evict)
- Caching disabled - used to keep data that logically belongs on the disk in memory to improve performance

# TLB

- **Observation:** Most programs make a large number of references to **a small number of pages**.
- **Solution:** Equip computers with a small hardware device, called **Translation Look-aside Buffer (TLB)** or **associative memory**, to map virtual addresses to physical addresses without using the page table.




# TLB – Translation Lookaside Buffer

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

A TLB to speed up paging

# Paging Hardware With TLB





# Inverted Page Table

- Usually, each process has a page table associated with it. One of drawbacks of this method is that each page table may consist of millions of entries.
- To solve this problem, an **inverted page table** can be used. There is one entry for each real page (frame) of memory.
- Each entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page.

# Inverted Page Table

- To illustrate this method, a simplified version of the implementation of the inverted page is described as:

**<process-id, page-number, offset>**

- Each entry: <process-id, page-number>.

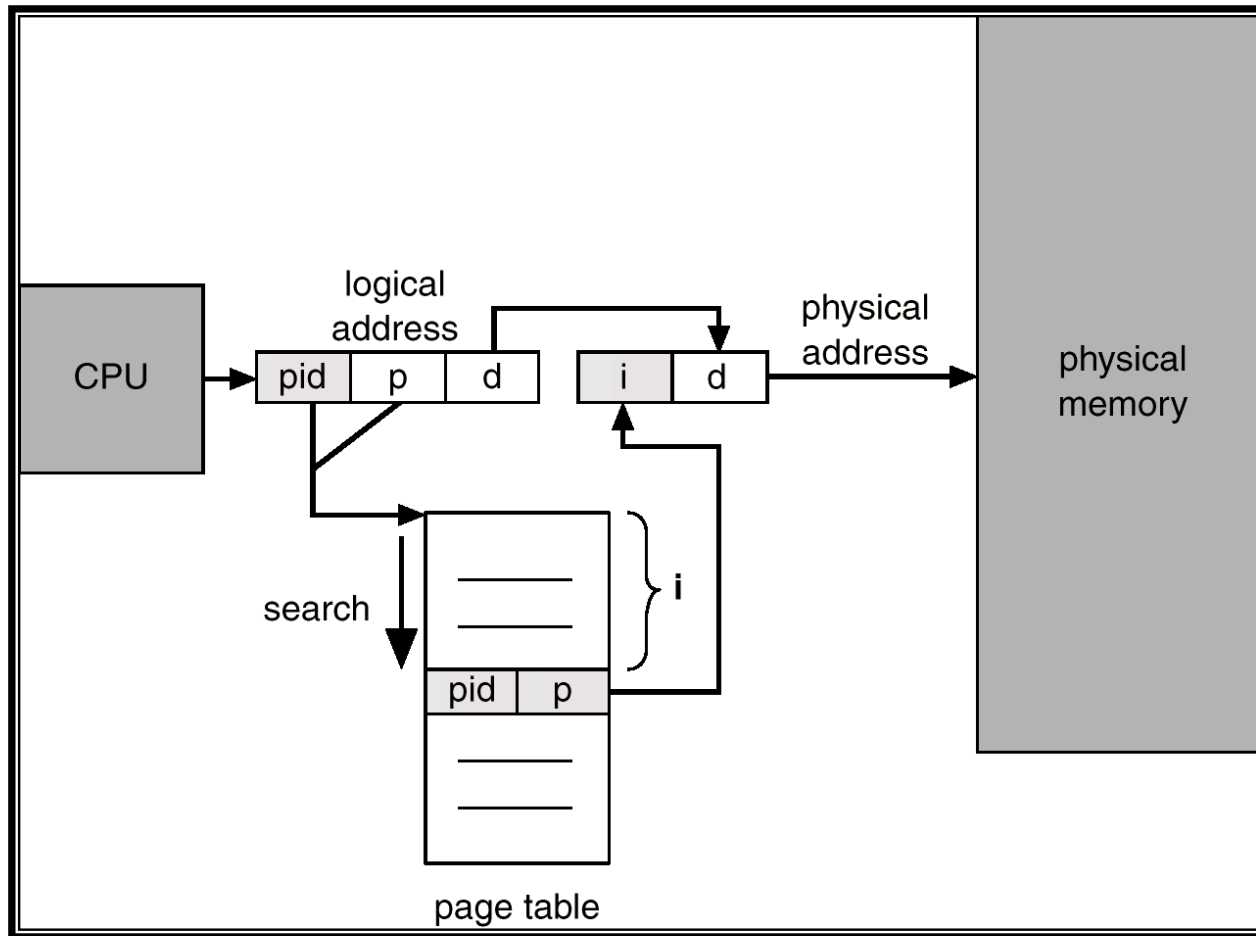
The inverted page table is then searched for a match. If a match  $i$  is found, then the physical address < $i$ , offset> is generated.

Otherwise, an illegal address access has been attempted.

- Although it decreases memory needed to store each page table, but **increases time needed to search the table** when a page reference occurs.



# Inverted Page Table Architecture



# Check Points

- What is the drawback of the bitmap method for free memory management?
- What is the purpose of virtual memory?
- What is page table?
- What is the purpose of using multi-level page table?
- What is the purpose of using TLB?
- What is invert page table?

# Presentation & Poster

- ① Group 1,10: Process Scheduling
- ② Group 2,11: The Banker's Algorithm
- ③ Group 3,12: Virtual Memory Paging
- ④ Group 4,13: Page Replacement Algorithm
- ⑤ Group 5,14: Memory Management (Link List)
- ⑥ Group 6,15: Disk space Management
- ⑦ Group 7,16: Interrupt/DMA
- ⑧ Group 8,17: Disk Arm Scheduling Algorithm
- ⑨ Group 9,18: Files: Link List Allocation



# 华南理工大学2019~2020学年度第一学期校历

年 月、 周次 星期	2019年																		2020年									
	8月	9月				10月				11月				12月				1月				2月						
	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	寒 假							
	星 期 一	星 期 二	星 期 三	星 期 四	星 期 五	星 期 六	星 期 日	星 期 一	星 期 二	星 期 三	星 期 四	星 期 五	星 期 六	星 期 日	星 期 一	星 期 二	星 期 三	星 期 四	星 期 五	星 期 六	星 期 日	星 期 一	星 期 二	星 期 三	星 期 四	星 期 五	星 期 六	星 期 日
	26	2	9	16	23	30	7	14	21	28	4	11	18	25	2	9	16	23	30	6	13	20	27	3	10	17		
	27	3	10	17	24	1	8	15	22	29	5	12	19	26	3	10	17	24	31	7	14	21	28	4	11	18		
	28	4	11	18	25	2	9	16	23	30	6	13	20	27	4	11	18	25	1	8	15	22	29	5	12	19		
	29	5	12	19	26	3	10	17	24	31	7	14	21	28	5	12	19	26	2	9	16	23	30	6	13	20		
	30	6	13	20	27	4	11	18	25	1	8	15	22	29	6	13	20	27	3	10	17	24	31	7	14	21		
	31	7	14	21	28	5	12	19	26	2	9	16	23	30	7	14	21	28	4	11	18	25	1	8	15	22		
	1	8	15	22	29	6	13	20	27	3	10	17	24	1	8	15	22	29	5	12	19	26	2	9	16	23		
	教 学 18 周																		考试2周	寒 假 6 周								

第6周	期一	5-8节	B3	138
第12周	星期二	5-8节	B3	231
第13周	星期一	1-4节	B3	138
第13周	星期三	1-4节	B3	138
第13周	星期五	1-4节	B3	234

